

1 Math

$$\sigma^s = \mathbb{E}[X^2] - \mathbb{E}[X]^2 \quad (1)$$

$$\frac{\partial}{\partial q_k} \text{soft}(q)_i = \text{soft}(q)_i (\delta_{i,k} - \text{soft}(q)_k) \quad (2)$$

$$KL(p||q) = \int p(x) \log\left(\frac{p(x)}{q(x)}\right) dx \quad (3)$$

$$(4)$$

2 NN and Optimization

RMSProp and **ADAM** are actly worse for simple landscapes. **Adam**:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \wedge v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \wedge \hat{m}_t = \frac{m_t}{1 - \beta_1^t} \wedge \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

init correction for sec momentum necessary (RMSProp+Mom is worse). **RMSProp**:

$$r_t = \alpha r_{t-1} + (1 - \alpha) g_t^2 \wedge u_t = -\frac{\eta}{\sqrt{r_t + \epsilon}} g^t \wedge w_{t+1} = w_t + \eta_t u_t$$

2ND-Order Optim $w_{t+1} = w_t - H_{\mathcal{L}}^{-1} \eta_t g_t$ weight updates by hessian but to big to compute :(.

Activation functions around 0, better not saturating, not bounded, center of it should be mean of inputs

Batchnorm whitening for activation functions, regularizes inference. Have linear function after batchnorm: at beginning its centered but can unlearn this.

3 RNN

$$c_t = \tanh(x_{t-1}) + Ux_t + b, \quad \mathcal{L} = \sum_t \mathcal{L}_t(c_t). \text{ Grad is chain Jacobs: } \frac{\partial c_t}{\partial c_k} = \prod_{j=k+1}^t \frac{\partial c_j}{\partial c_{j-1}} \text{ and}$$

$$\frac{\partial \mathcal{L}}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_t}{\partial c_t} \frac{\partial c_t}{\partial c_{\tau}} \frac{\partial c_{\tau}}{\partial W} \text{ with restr:}$$

$$\|\frac{\partial c_{t+1}}{\partial c_t}\| \leq \|W^T\| \cdot \|\text{diag}(\sigma'(c_t))\|$$

$$\text{If } \|\frac{\partial c_k}{\partial c_{k-1}}\| \leq \frac{1}{\lambda_{max}} \|\text{diag}(\max(\sigma'(\cdot)))\| < 1 \text{ then}$$

$$\pi_{k=1}^{\tau} \frac{\partial c_k}{\partial c_{k-1}} \text{ goes zero exp, van grads. Opposite is expld grads.}$$

4 GNN

DeepWalk: randomwalk + LSTM with skip-gram, works not good new nodes need retrain.

5 Deep Generative Models

Boltzman dist: $p(x) = \frac{1}{Z} \exp(-E(x))$. Comp. of normal. Const. Z difficult. **Boltzmann machine** $E(x) = -x^T W x - b^T x$ x is 256² big. Instead **RBM** $E(x) = -x^T W h - b^T x - c^T h$ with latent h.

5.1 GAN

implicit density, sampling from PDF $\mathcal{L} = -12 \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_{z \sim p_z} \log(1 - D(G(z)))$ For better learning train for G the opposite, for approx ML estimate: $J^G = -\frac{1}{2} \mathbb{E}_z \exp(\sigma^{-1}(D(G(z))))$ (one opt, Goodfellow). Normal object resembles minimizing Jesnon-Shannon divergence: $D_{JS}(a||b) = 0.5 D_{KL}(a|| (a+b)/2) + 0.5 D_{KL}(b|| (a+b)/2)$ GAN problems: minimax instability, van grads, mode collapse (not due to divergence probably) Improvements: Wasserstein dist, cBN, cGAN, label smoothing

5.2 Variational Inference

How est. posterior: MCMC or var. infer.: $\phi^* = \arg \min_{\phi} KL(q(\theta|\phi)||p(\theta|x))$, rev divergence (underestimate var, overest. with forward). **ELBO** $\mathbb{E}_{q_{\phi}(\theta)}[\log p(x|\theta)] - KL(q_{\phi}(\theta)||p(\theta)) = \mathbb{E}_{q_{\phi}(\theta)}[\log p(x|\theta)] + \mathbb{E}_{q_{\phi}(\theta)}[\log p(\theta)] - \mathbb{E}_{q_{\phi}(\theta)}[\log q(\theta)]$ with that $\log p(x) = ELBO_{\theta, \phi}(x) + KL(q_{\phi}(\theta)||p(\theta|x))$. **ELBO** is vari. free Enrgy. Backprop in VAE: use **REINFORCE** to approx grad (high var grads slow down) or reparam trick

5.3 Normalizing Flows

$\log p(x) = \log \pi_0(z_0 - \sum_i^K | \det \frac{df_i}{dz_{i-1}} |)$ requirements: f_i must be easily invertible and the Jacobian must be computable

6 Bayesian Deep Learning

Benefits of Bayesian: ensemble makes better accuracies, uncertainty estimates, sparsity makes model compression, active learning, distributed learning. **Epistemic uncertainty** ignorance which model generated the data. More data reduces this. For safety critical stuff, small datasets. **Aleatoric uncertainty** ignorance about

the nature of the data. *Heteroscedastic* uncertainty about specific data $\mathcal{L} = \frac{\|y_i - \hat{y}_i\|^2}{\sigma_i^2} + \log \sigma_i$, *homoscedastic* uncertainty about the task, we might reduce by combining tasks. \mathcal{L} same but without idx. **MC Dropout** have d. during inference (by Bernoulli as vari. dist.) Then model prec. $\tau = \frac{l^2 p}{2N\lambda}$.

7 Deep Sequential models

7.1 Autoregressive models

With sequential data we have:

$x = [x_1, \dots, x_k] \implies p(x) = \prod_{k=1}^D p(x_k | x_{j < k})$ thus no param sharing and no ∞ chains $\implies p(x)$ is tractable.

NADE: fixed masks, conditionals modeled as MoG. **MADE**: masked conv on an autoencoder. **PixelRNN** seq. order over rows and channel R,G and B. Conditionals modeled with LSTM. Slow train and gen, but good gen. **PixelCNN** model conds with masked convs. Is worse than RNN cause blind spot. Fix by having convs for left row and everything above cascading. Output 8bit softmax **GatedPixelCNN** use two conv stacks, horiz and vert to not have blind sport **PixelCNN++** dropout/whole pixels/discr log mix likelihood (from continuos output). **PixelCNN** is too powerful **PixelVAE** VAE+PixelCNN as the networks

8 DeepRL

Goal: max fut rewards (Q-func, value).

$$Q^{\pi}(s_t, a_t) = \mathbb{E}(r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s_t, a_t) = \mathbb{E}_{s', a'}(r + \gamma Q^{\pi}(s', a') | s_t, a_t) \text{ (Bellman eq).}$$

Approaches, -based, policy, value, model.

Optimal Value Function

$$Q^*(s, a) = t_{t+1} + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) = \mathbb{E}_{s'}(t + \gamma \max_{a'} Q^*(s', a') | s, a)$$

Value based: **Q-learning**: minimize

$\min(r + \gamma \max_{a'} Q_t(s', a') - Q_t(s, a))^2$ $Q_t(s, a)$ is your prev est Q-Value for state s and action a. The other stuff is your new estimate at new state, action. You want to minimize this to get better.

Gradient ist then $\frac{\partial \mathcal{L}}{\partial \theta} =$

$$\mathbb{E}[-2 \cdot (r + \gamma \max_{a'} Q(s', a', \theta) - Q(s, a, \theta)) \frac{\partial Q(s, a, \theta)}{\partial \theta}]$$

Is unstable because target depends on Q, also seq breaks independence assumpt, highly correlated samples break SGD. Solut1: exp replay, play random steps from other history. Solt2: have a sec network which is updated once in a while to calc

targets so that dies not interfere with grad calc stability. *Other tricks*: clip rewards to -1,1, skip frames

Policy Optimization: q-func often too expensive, must account for all states/actions. Instead directly learn policy $\pi_{\theta}(a|s)$:

$$\frac{\partial \mathcal{L}}{\partial w} = \mathbb{E}[\frac{\partial \log \pi(a|s, w)}{\partial w} Q^{\pi}(s, a)] \text{ (deterministic) or}$$

$$\frac{\partial \mathcal{L}}{\partial w} = \mathbb{E}[\frac{\partial Q^{\pi}(s, a)}{\partial a} \frac{\partial a}{\partial w}] \text{ (stochastic } a = \pi(s)) \text{ compute}$$

$$\nabla_{\theta} \log p(x; \theta) = \frac{\nabla_{\theta} p(x; \theta)}{p(x; \theta)}$$